

Avalon-MM SDRAM Controller IP Core

Block Diagrams and Descriptions


IP core	<code>altera_avalon_mm_sdram_controller</code>
Core version	1.0
Document version	1.0
Date	August 2026
Target	Intel FPGA / Quartus Prime, Platform Designer

Contents

1. System context	3
2. One open row, or one per bank	6
3. Internal architecture	9
4. Address map	12
Revision history	14

1. System context

System context



The same three interfaces the SDRAM Controller Intel FPGA IP presents: a word-addressed Avalon-MM memory slave named `s1` carrying the legacy `az_/za_` signals, a conduit named `wire` carrying the SDRAM pins, and a clock and reset sink. Swapping one component for the other in an existing Platform Designer system leaves every connection and every address assignment alone.

Figure 1. System context

The controller sits between an Avalon-MM master and an SDR SDRAM device. It presents three interfaces, and they are deliberately the same three the SDRAM Controller Intel FPGA IP presents:

Interface	Kind	What it carries
s1	Avalon-MM slave, word-addressed	az_addr, az_be_n, az_cs, az_data, az_rd_n, az_wr_n, za_data, za_valid, za_waitrequest
wire	conduit	The SDRAM pins: zs_addr, zs_ba, zs_cas_n, zs_cke, zs_cs_n, zs_dq, zs_dqm, zs_ras_n, zs_we_n
clk / reset	clock and reset sink	One clock domain. Nothing here crosses domains

Because the interfaces match, replacing the Intel component with this one in an existing Platform Designer system leaves every connection and every address assignment alone.


Note

The clock rate is not a parameter. Platform Designer takes it from the clock source the component is connected to and passes it to the HDL, which converts every device timing against it. A controller configured for 100 MHz and clocked at 143 is not a mistake that can be made here.

2. One open row, or one per bank


One open row, or one per bank

Single open row



An access to any other bank must close the open row and open a new one: PRECHARGE, tRP, ACTIVATE, tRCD. So must a read after a write, because the fast path also requires the direction to match.

One open row per bank



Four rows stay open at once. An access to another bank needs no row command at all, and a read/write turnaround inside an open row costs 0 cycles write-to-read and CAS+1 read-to-write, with no row command either way.

This is the whole design, and it is measurable rather than arguable. Four banks with one row open in each, accessed in rotation, needs exactly four ACTIVATEs and no PRECHARGE - which is what the testbench asserts, and what takes that access pattern from 21.9 MB/s to 195.5 MB/s.

Figure 2. Row tracking

This is the entire design, and the reason the core exists.

An SDRAM bank has one row buffer. Reading or writing a location requires that location's row to be **open** — ACTIVATED into the buffer. Opening a different row in the same bank means PRECHARGE, then tRP, then ACTIVATE, then tRCD; on the DE10-Lite's part at 100 MHz that is $2 + 2 + 2 = 6$ dead cycles before the next access can even be issued.

A device has **four** banks and therefore four row buffers. A controller may keep a row open in every one of them at the same time.

The controller this replaces keeps one. Its fast path is:

```
assign pending = csn_match && rnw_match && bank_match && row_match && !f_empty;
```

so an access falls out of it — into a full row cycle — if the bank differs, if the row differs, **or if the direction differs**. That last condition is the expensive one: a read following a write inside a single open row, which the device can do back to back, costs a full PRECHARGE/ACTIVATE pair.


This core tracks `row_open[4]` and `open_row[4]`, and treats a turnaround as the datasheet describes it: 0 cycles write → read, CAS+1 read → write, no row command either way.

Caution

A per-bank design needs two timing parameters a single-row design never encounters. **tRAS** bounds how soon a row may be closed after opening, and **tRRD** bounds how soon a second bank may be activated after the first. Neither can be violated by a controller that only ever has one row open, and both can be by this one. They are checked in simulation by `tb/sdram_timing_check.sv`.

3. Internal architecture

Internal architecture



Every SDRAM output is registered, so a command decided in one cycle is sampled by the device in the next. The timing counters are loaded accordingly; getting that off by one is either a wasted cycle on every row change, or a violation.

Figure 3. Internal architecture

Command FIFO. Accepted transfers are buffered `FIFO_DEPTH` deep. The buffer does two jobs: it decouples the master from row-change stalls, and it gives the scheduler something to look at beyond the access it is currently serving.

Scheduler. One decision per cycle, in priority order:

Condition	Command issued
The head's row is open, tRCD met, and the bus is free in that direction	READ or WRITE, and the entry is retired
The head's bank is open on the wrong row, and tRAS/tWR permit	PRECHARGE
The head's bank is closed, and tRP/tRC/tRRD permit	ACTIVATE
None of the above, and LOOKAHEAD is on	PRECHARGE or ACTIVATE for the <i>next</i> entry, if it is in a different bank

Bank state. `row_open[]` and `open_row[]` are the design's central claim. If they ever disagree with the commands actually issued, every subsequent access lands somewhere else — silently. The testbench maintains an independent shadow of this state, rebuilt from the command bus, and asserts the two match.

Timing gates. A down-counter per bank for tRCD, tRAS, tRP, tRC and tWR, and global counters for tRRD, tRFC and the read/write bus turnaround. Every one is derived from a picosecond parameter and the clock rate, inside the HDL.

Refresh. A timer accumulates outstanding refreshes. They are spent when the buffer is empty — where they cost nothing — or forced through once `REF_MAX_PEND` have accumulated. JEDEC permits up to 8 postponed.

Read return. A shift register `CAS_LAT` deep carries a marker; the data itself is captured straight off the bus. No address is stored, because SDRAM returns read data strictly in order.

Note

Every SDRAM output is registered. A command decided in cycle k is sampled by the device on edge $k+1$, so the timing counters are loaded with $(\text{cycles} - 1)$. Getting that off by one is either a wasted cycle on every row change, or a violation — and it is invisible in a functional simulation.

4. Address map

Address map

ADDR_MAP 0 - compatible (default). The map the SDRAM Controller Intel FPGA IP uses.



The bank bits are SPLIT: bank[0] sits directly above the column, so a purely ascending address alternates bank every 1024 words and advances the row every 2048. Even a perfectly sequential stream changes bank constantly - which a one-open-row design pays for and this one does not.

ADDR_MAP 1 - conventional. New designs only.



Selecting the conventional map moves every address in memory. It is the right choice for a new design and the wrong one for replacing the Intel core in a system that already exists, which is why the component's validation warns about it at generation time.

Figure 4. Address map

ADDR_MAP 0, the default, reproduces the layout the Intel core uses:

Field	Bits, at the default geometry
bank[1]	addr[24]
row	addr[23:11]
bank[0]	addr[10]
column	addr[9:0]

The bank bits are **split**. bank[0] sits directly above the column, so an ascending address alternates bank every 1024 words and advances the row every 2048. Even a perfectly sequential stream changes bank constantly — which is precisely the traffic a one-open-row design handles badly and a per-bank design does not notice.

ADDR_MAP 1 is the conventional {row, bank, column}.

Caution

Selecting the conventional map moves every address in memory. It is the right choice for a new design and the wrong one for replacing the Intel core in a system that already exists. The Platform Designer component emits a warning when it is selected.

Revision history

Document version	Core version	Date	Change
1.0	1.0	August 2026	First release